



**NANYANG
TECHNOLOGICAL
UNIVERSITY**

CE1007/CZ1007 DATA STRUCTURES

Lecture 09: Binary Search Trees


Dr. Owen Noel Newton Fernando

College of Engineering

School of Computer Science and Engineering


RECALL: WHY TREES?

- Model layouts with hierarchical relationships between items
 - Chain of command in the army
 - Personnel structure in a company
 - (Binary tree structure is limited because each node can have at most two children)
- Tree structures also allow us to
 - Some problems require a tree structure: some games, most optimization problems, etc.
 - Allow us to do the following very quickly: (we'll see that in the following lectures)
 - **Search for a node with a given value**
 - **Add a given value to a list**
 - **Delete a given value from a list**



- **Item Search**


- Binary Search Trees (BST)
- BST Operations:
 - Traversal
 - Inserting a node
 - Removing a node



ITEM SEARCH-LINKED LIST

inefficient

Given a linked list of names, how do we check whether a given name(e.g., **Irit**) is in the list?



```
while (cur!=NULL) {  
    if cur->item == "Irit"  
        found and stop searching;  
    else  
        cur = cur->next; }
```


How many nodes are visited during search?
--best case: 1 node (John)
--worst case: 7 nodes (Simon)
--avg. case: $(1+2+3+...+7)/7=4$ nodes

ITEM SEARCH-LINKED LIST

inefficient

Given a binary tree of names, how do we check whether a given name(e.g., **Brian**) is in the tree?

Use the TreeTraversal (Pre-order) template, to check every node



```
TreeTraversal(Node N)
  If N==NULL return;
  if N.item=given_name return;
  TreeTraversal(LeftChild);
  TreeTraversal(RightChild);
  Return;
```

How many nodes are visited during search?


--best case: 1 node (Anna)

--worst case: 7 nodes (Jane)


--Avg. case: $(1+2+3+...+7)/7=4$ nodes

IS THERE A WAY TO ORGANIZE THE DATA SO THAT ITEM SEARCH CAN BE MORE EFFICIENT?

- What's a good way to arrange data in our tree so that we can efficiently store and retrieve items?
- How do we efficiently (minimize number of operations) find an item in a binary tree?
- Given a list of items, how do we check if some given item is contained inside?



For the list:




Sort it using alphabetical order:

Divide them into groups

Search the
given name X


- Pick the one in the middle, "Jane", if X!="Jane":
 - For names before "Jane", lookup its left subgroup, ignore its right subgroup
 - Pick the one in the middle of the subgroup, "Brian", if X!="Brian":
 - names before "Brian", lookup its left subgroup
 - names after "Brian", lookup its right subgroup
 - For names after "Jane", lookup its right subgroup, ignore its left subgroup
 - Pick the one in the middle of the subgroup, "Peter", if X!="Peter":
 - <"Peter", lookup its left subgroup
 - >"Peter", lookup its right subgroup



IT FORMS A BINARY SEARCH TREE (BST)

Divide them into groups


- Pick the one in the middle, "Jane", if $X \neq \text{"Jane"}$:
 - For names before "Jane", lookup its left subgroup, ignore its right subgroup
 - Pick the one in the middle of the subgroup, "Brian", if $X \neq \text{"Brian"}$:
 - names before "Brian", lookup its left subgroup
 - names after "Brian", lookup its right subgroup
 - For names after "Jane", lookup its right subgroup, ignore its left subgroup
 - Pick the one in the middle of the subgroup, "Peter", if $X \neq \text{"Peter"}$:
 - $< \text{"Peter"}$, lookup its left subgroup
 - $> \text{"Peter"}$, lookup its right subgroup



- Item Search

- **Binary Search Trees (BST)**

- BST Operations:
 - Traversal
 - Inserting a node
 - Removing a node



BINARY SEARCH TREE(BST)


- BSTs are a special form of BT

- **BST rule:**

At every node **C**,

L < **C** < **R**, where

- **C** is the data in the current node
- **L** represents the data in any/ all nodes from C's left subtree
- **R** represents the data in any/all nodes from C's right subtree




BINARY SEARCH TREE

- BSTs are a special form of BT
- At every node C,
 $L \leq C \leq R$, where

- C is the data in the current node
- L represents the data in the left child
- R represents the data in the right child


NO = in the BST!
There must be no
duplicate nodes in
BST!



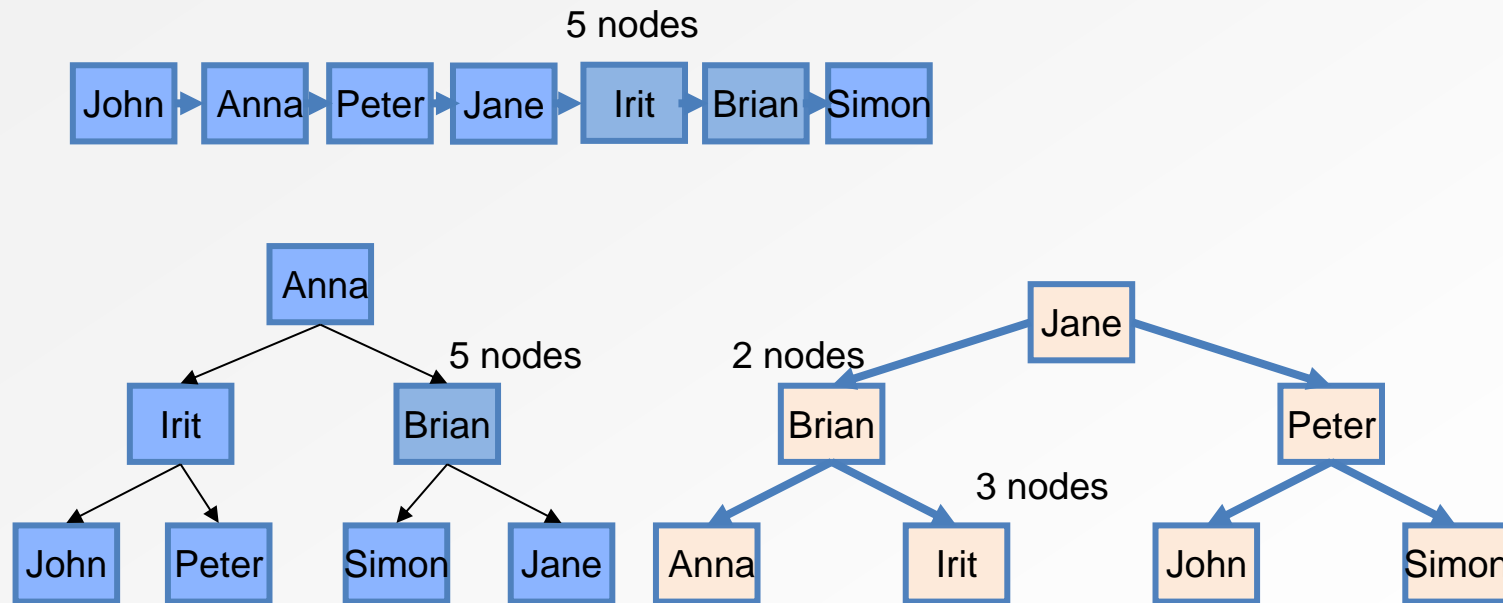
This is not a BST!

BST IS FOR EFFICIENT ITEM SEARCH

- This is a BST, satisfies ' $L < C < R$ '



BST IS EFFICIENT FOR ITEM SEARCH



How many nodes are visited during search?

--best case: 1 node (Anna)

--worst case: 7 nodes (Jane)

--Avg. case: $(1+2+3+\dots+7)/7=4$ nodes

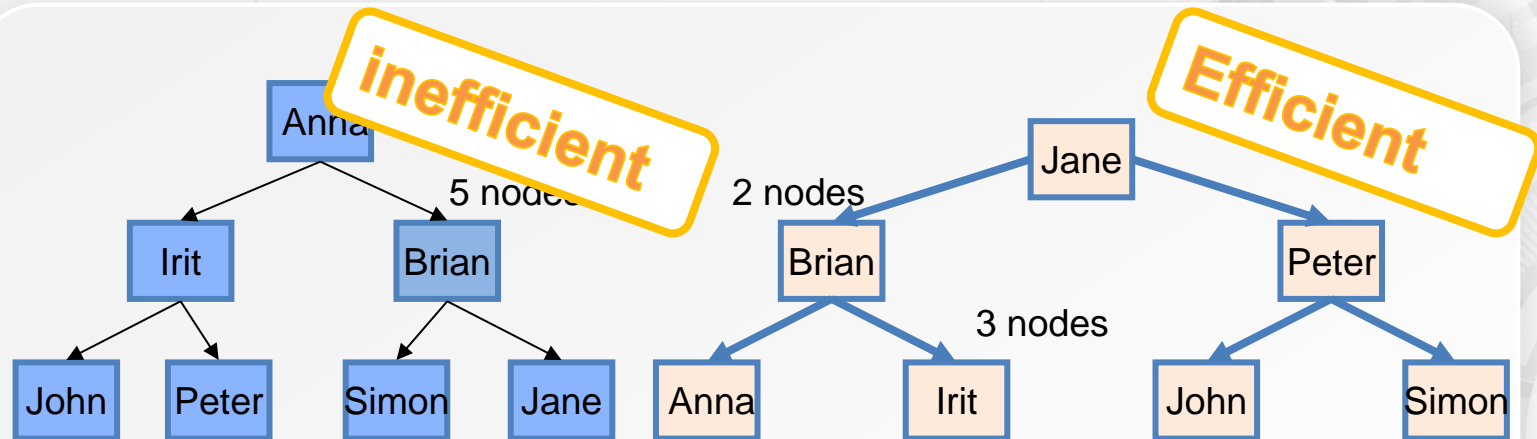
How many nodes are visited during search?

--best case: 1 node (Jane)

--worst case: 3 nodes (Anna)

--Avg. case: $(1+2*2+3*4)/7=2.43$ nodes

BST IS EFFICIENT FOR ITEM SEARCH



How many nodes are visited during search?

In general, for a BT with n nodes:

--best case: First node in traversal

--worst case: **Last node in traversal, n**

How many nodes are visited during search?

In general, for a BST with n nodes:

--best case: First node in traversal

--worst case:

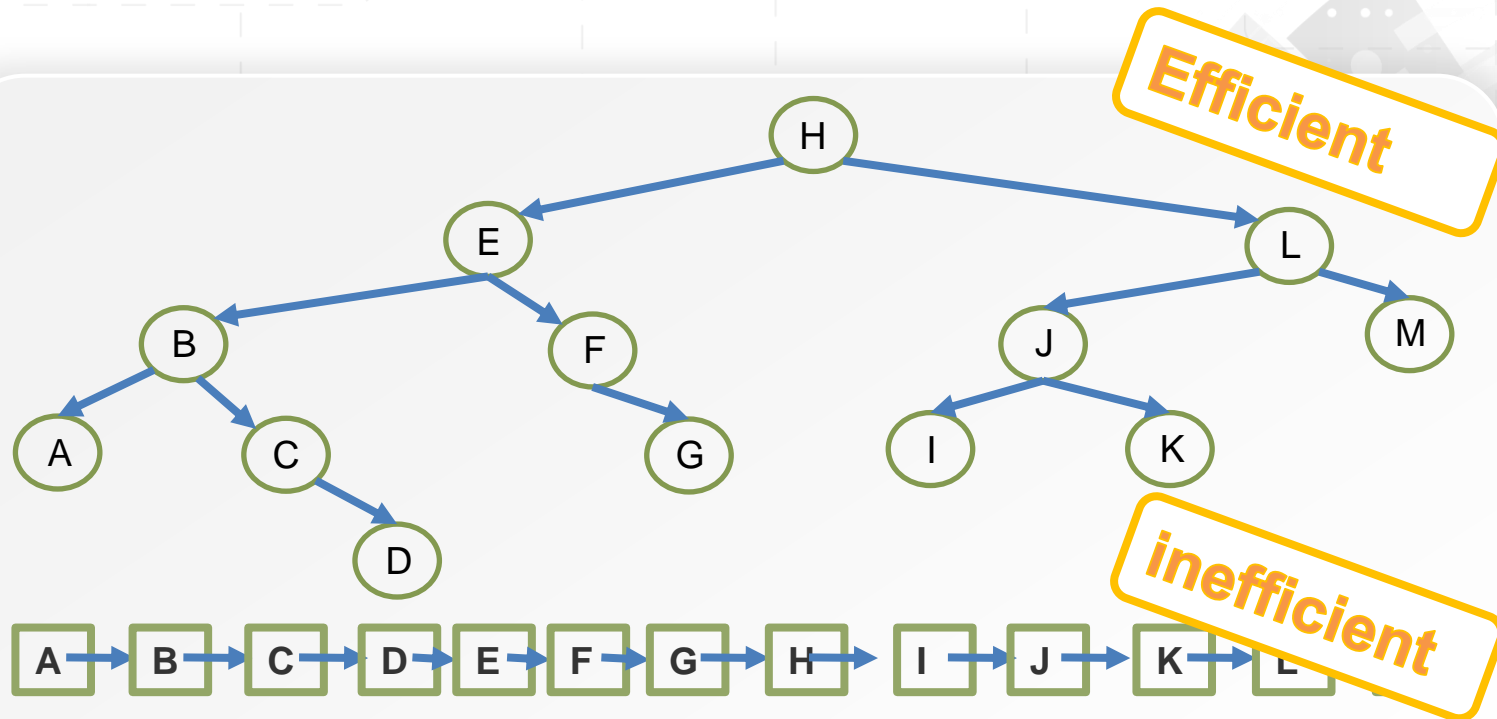
leaf node: the height of the root + 1

Minimal height $H = \lfloor \log_2 n \rfloor$

As n becomes a big number, the difference between them becomes even greater

Height of a node = number of links from that node to the deepest leaf node



ANOTHER BST



- How do we check if a given item X is stored in the tree?
- To find D, we need to visit H-E-B-C-D, 5 nodes.
- To find any node in the tree, at most visit 5 nodes.
- But for the linked list, the worst case need to visit all nodes - 13 nodes.

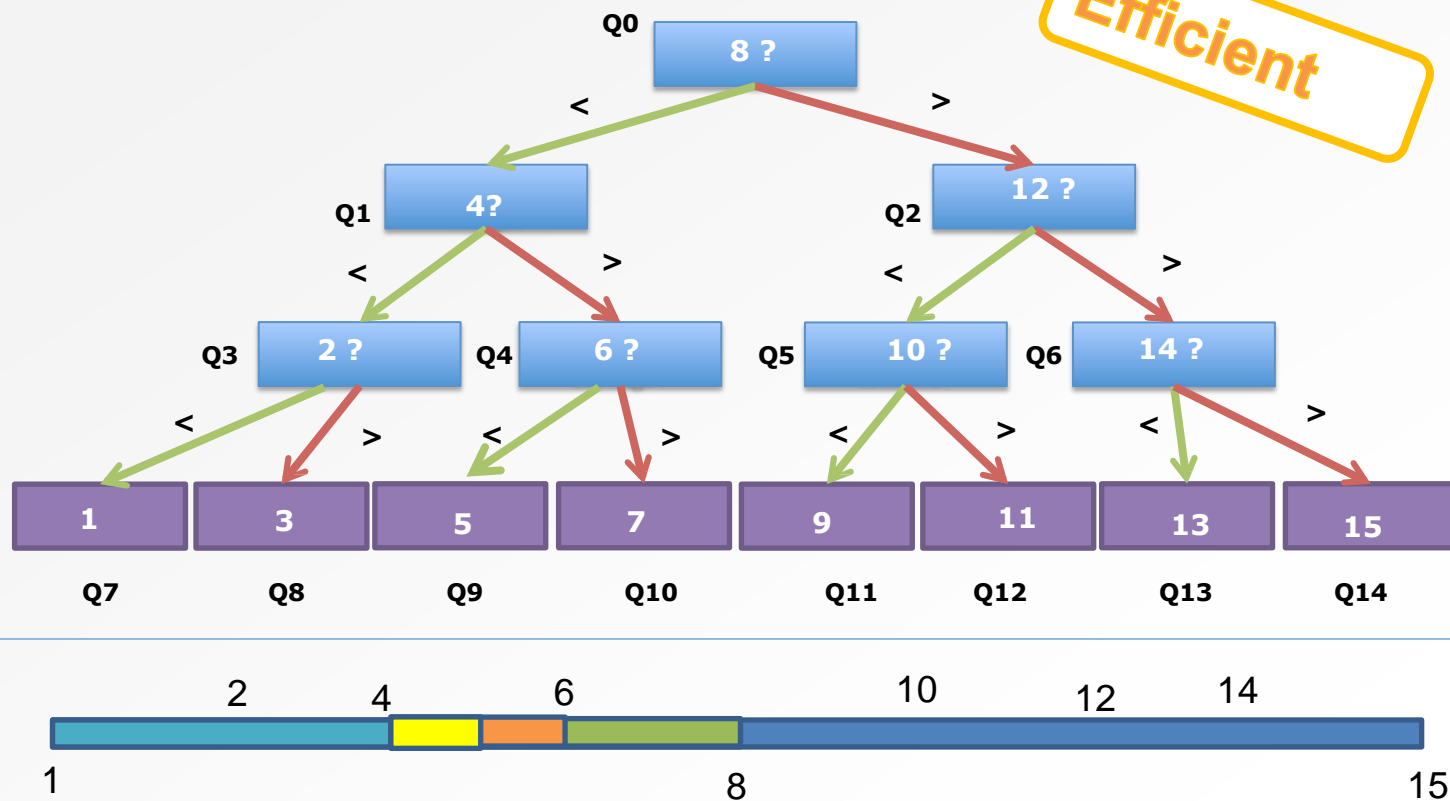
NOTICE: NOT ALL BST ARE EFFICIENT FOR SEARCH

- What does a good/bad BST look like?
- Two possible BST representations of the list



RECALL: NUMBER GAME [1, 15]


- How many questions do you ask to guess the number?
- Best case: 1 question
- Worst case: 4 questions




RECALL: NUMBER GAME [1, 15]

- How many questions do you ask to guess the number?
- If the strategy is to guess **1 first, and then 2, 3, ..., 15**
- Best case: **1** question
- Worst case: **15** questions



inefficient



EXERCISE: WHICH BINARY TREES IS EFFICIENT FOR SEARCH ?



THINK: WHAT IS THE VISITING SEQUENCE USING IN-ORDER TRAVERSAL FOR A BST?




Output:

A B C D E F G H I

?


MAPPING: TREE(IN-ORDER) → LIST



- If we draw the BST carefully:
 - **Left subtree on the left side of the current node;**
 - **Right subtree on the right side of the current node;**
- Mapping to X-axis will produce **a sorted list.**


FEATURES

- $L < C < R$ rule ensures sorted order
- BST's in-order traversal produces a sorted list!




FEATURES

- The binary-search-tree property guarantees that:
 - The **minimum** is located at the **left-most** node
 - The **maximum** is located at the **right-most** node



- Item Search
- Binary Search Trees (BST)
- BST Operations:
 - **Traversal**
 - Inserting a node
 - Removing a node



BINARY SEARCH TREE(BST)


- BSTs are a special form of BT

- **BST rule:**

At every node **C**,


L < **C** < **R**, where

- **C** is the data in the current node
- **L** represents the data in any/ all nodes from C's left subtree
- **R** represents the data in any/all nodes from C's right subtree



BST TRAVERSAL (BSTT)

- BSTT() traverses a BST to search for a node with a matching item
- Begin with TreeTraversal template



```
void BSTT(BTNode *cur, char c){  
    if (cur == NULL)  
        return;  
  
    // Do something  
  
    BSTT(cur->left);  
    BSTT(cur->right);  
}
```


Do something with the
current node's data

Visit the left child node

Visit the right child node

BST TRAVERSAL (BSTT)

- Now, at each node, we need to determine which subtree to keep visiting (and which subtree to ignore)



```
void BSTT(BTNode *cur, char c){  
    if (cur == NULL) return;  
    if (c==cur->item)  
    { printf("found!\n"); return;}  
    if (c < cur->item)  
        BSTT(cur->left,c);  
    else  
        BSTT(cur->right,c);  
}
```


Do something with the current node's data

Visit the left child node

Visit the right child node

BST TRAVERSAL (BSTT)


- Check the traversal pattern for **BSTT(root, 'B')**



```
void BSTT(BTNode *cur, char c){  
    if (cur == NULL) return;  
    if (c==cur->item)  
    { printf("found!\n"); return;}  
    if (c < cur->item)  
        BSTT(cur->left,c);  
    else  
        BSTT(cur->right,c);  
}
```

BST TRAVERSAL (BSTT)


- Check the traversal pattern for **BSTT(root, 'K')**



```
void BSTT(BTNode *cur, char c){  
    if (cur == NULL) return;  
    if (c==cur->item)  
    { printf("found!\n"); return;}  
    if (c < cur->item)  
        BSTT(cur->left,c);  
    else  
        BSTT(cur->right,c);  
}
```


BST TRAVERSAL (BSTT)

- What if the item doesn't exist?
- If we remove node 'D', and then check the traversal pattern for **BSTT(root, 'D')**




```
void BSTT(BTNode *cur, char c){  
    if (cur == NULL) {printf("can't find!");return; }  
    if (c==cur->item)  
    { printf("found!\n"); return;}  
    if (c < cur->item)  
        BSTT(cur->left,c);  
    else  
        BSTT(cur->right,c);  
}
```

- Item Search
- Binary Search Trees (BST)
- BST Operations:
 - Traversal
 - **Inserting a node**
 - Removing a node




INSERTING A NODE INTO A BST

- Given an existing BST, an insertion operation must result in a BST
- How do we know where to place a new node 'D'?




INSERTING A NODE INTO A BST

- Key point:
 - Given an existing BST and a new value to store, there is always a unique position for the new value




INSERTING A NODE INTO A BST

1. Use BSTT() to get to the correct empty location
2. Add the new node




INSERTING A NODE INTO A BST

1. Use $\text{BSTT}(\text{root}, 'D')$ to get to the correct empty location to insert 'D'
2. Add the new node 'D'




INSERTING A NODE INTO A BST

- Node insertion is relatively simple!
- Further exercise: Try Inserting 'Z'



- Item Search
- Binary Search Trees (BST)
- BST Operations:
 - Traversal
 - Inserting a node
 - **Removing a node**

After removal, the tree is still a BST




REMOVING A NODE FROM A BST

- Node removal is more complicated
- Beginning with a BST, the resulting tree after removing a node must still be a BST

Obey the BST rule: $L < C < R$


REMOVING A NODE FROM A BST

- Remove node X - a bit tricky
- 3 cases:
 1. x has no children:
 - Remove x
 2. x has one child y:
 - Replace x with y
 3. x has two children:
 - Swap x with successor
 - Perform case 1 or 2 to remove it




REMOVING A NODE FROM A BST

- Remove node X - a bit tricky
- 3 cases:
 - 1. x has no children:**
 - **Remove x**
 2. x has one child y:
 - Replace x with y
 3. x has two children:
 - Swap x with successor
 - Perform case 1 or 2 to remove it



REMOVING A NODE FROM A BST

- Remove node X - a bit tricky
- 3 cases:
 1. x has no children:
 - Remove x
 2. **x has one child y:**
 - **Replace x with y**
 3. x has two children:
 - Swap x with successor
 - Perform case 1 or 2 to remove it



REMOVING A NODE FROM A BST

- Remove node X - a bit tricky

- 3 cases:

1. x has no children:


- Remove x

2. x has one child y:

- Replace x with y

3. x has two children:

- **Swap x with successor**
- **Perform case 1 or 2 to remove it**



WHAT IS THE SUCCESSOR OF X?


Replacing a node with its in-order successor ensures that the BST rule ($L < C < R$) is maintained

In-order traversal of a BST produce a sorted list (in ascending order)

Successor is:

- The node immediately after it in the sorted list, or
- The next node visited using an in-order traversal

X has two children, so X's successor is minimum node in its right subtree.
E.g.: H's successor is I, E's successor is F, J's successor is K.



REMOVING A NODE FROM A BST

- Remove node X - a bit tricky

- 3 cases:

1. x has no children:


- Remove x

2. x has one child y:

- Replace x with y

3. x has two children:

- **Swap x with successor**
- **Perform case 1 or 2 to remove it**



QUESTIONS

- **Why will case 3 always go to case 1 or case 2?**

A: because when X has 2 children, its successor is the minimum in its right subtree, so the successor should not have left child.

It might have no child(case 1) or one right child(case 2).

- **Could we swap x with predecessor instead of successor?**

A: yes.

REMOVING A NODE FROM A BST

- Remove node X - a bit tricky

- 3 cases:

1. x has no children:


- Remove x

2. x has one child y:

- Replace x with y

3. x has two children:

- **Swap x with successor**
- **Perform case 1 or 2 to remove it**



WHAT IS THE SUCCESSOR OF X?


Replacing a node with its in-order **predecessor** ensures that the BST rule ($L < C < R$) is maintained

In-order traversal of a BST produce a sorted list (in ascending order)

Successor/predecessor:

- The node immediately after/**before** it in the sorted list
- The next/**previous** node visited using an in-order traversal

X has two children, so X's predecessor is maximum node in its left subtree.
E.g.: H's predecessor is G, E's predecessor is D, J's predecessor is I.



TODAY YOU SHOULD BE ABLE TO

- Define a Binary Search Tree
- From a list, how do we construct a Binary Search Tree?
Is it efficient?
- How do we traverse a BST to search a item?
- How do we insert/remove a node from a BST?